

Article

Dual Machine Learning Framework for Predicting Long-Term Glycemic Change and Prediabetes Risk in Young Taiwanese Men

Chung-Chi Yang ^{1,2,3,4}, Sheng-Tang Wu ^{5,6}, Ta-Wei Chu ^{7,8}, Chi-Hao Liu ^{9,10,*} and Yung-Jen Chuang ^{3,4,*}

- ¹ Division of Cardiovascular Medicine, Taoyuan Armed Forces General Hospital, Taoyuan 325208, Taiwan; t220979@gmail.com
 - ² Cardiovascular Division, Tri-Service General Hospital, National Defense Medical University, Taipei 11490, Taiwan
 - ³ School of Medicine, National Tsing Hua University, Hsinchu 30044, Taiwan
 - ⁴ Institute of Bioinformatics and Structural Biology, National Tsing Hua University, Hsinchu 30044, Taiwan
 - ⁵ Division of Urology, Department of Surgery, Tri-Service General Hospital, National Defense Medical University, Taipei 11490, Taiwan; doc20283@gmail.com
 - ⁶ Division of Urology, Department of Surgery, Kaohsiung Armed Forces General Hospital, Kaohsiung 802301, Taiwan
 - ⁷ Department of Obstetrics and Gynecology, Tri-Service General Hospital, National Defense Medical University, Taipei 11490, Taiwan; taweichu@gmail.com
 - ⁸ MJ Health Research Foundation, Taipei 114066, Taiwan
 - ⁹ Division of Nephrology, Department of Internal Medicine, Kaohsiung Armed Forces General Hospital, Kaohsiung 802301, Taiwan
 - ¹⁰ School of Medicine, National Defense Medical University, Taipei 11490, Taiwan
- * Correspondence: colinliu1201@gmail.com (C.-H.L.); yjchuang@life.nthu.edu.tw (Y.-J.C.); Tel.: +886-7-749-6751 (C.-H.L.); +886-3-574-2764 (Y.-J.C.)

Abstract

Background: Early detection of dysglycemia in young adults is important but underexplored. This study aimed to (1) predict long-term changes in fasting plasma glucose (δ -FPG) and (2) classify future prediabetes using complementary machine learning (ML) approaches.

Methods: We analyzed 6247 Taiwanese men aged 18–35 years (mean follow-up 5.9 years). For δ -FPG (continuous outcome), random forest, stochastic gradient boosting (SGB), extreme gradient boosting (XGBoost), and elastic net were compared with multiple linear regression using Symmetric mean absolute percentage error (SMAPE), Root mean squared error (RMSE), Relative absolute error(RAE), and Root relative squared error (RRSE) Sensitivity analyses excluded baseline FPG (FPG_{base}). Shapley additive explanations(SHAP) values provided interpretability, and stability was assessed across 10 repeated train–test cycles with confidence intervals. For prediabetes (binary outcome), an XGBoost classifier was trained on top predictors, with class imbalance corrected by SMOTE-Tomek. Calibration and decision-curve analysis (DCA) were also performed. **Results:** ML models consistently outperformed regression on all error metrics. FPG_{base} was the dominant predictor in full models (100% importance). Without FPG_{base} , key predictors included body fat, white blood cell count, age, thyroid-stimulating hormone, triglycerides, and low-density lipoprotein cholesterol. The prediabetes classifier achieved accuracy 0.788, precision 0.791, sensitivity 0.995, ROC-AUC 0.667, and PR-AUC 0.873. At a high-sensitivity threshold (0.2892), sensitivity reached 99.53% (specificity 47.46%); at a balanced threshold (0.5683), sensitivity was 88.69% and specificity was 90.61%. Calibration was acceptable (Brier 0.1754), and DCA indicated clinical utility. **Conclusions:** FPG_{base} is the strongest predictor of glycemic change, but adiposity, inflammation, thyroid status, and lipids remain informative. A dual interpretable ML framework offers clinically actionable tools for screening and risk stratification in young men.



Academic Editor: Mara Carsote

Received: 8 August 2025

Revised: 28 September 2025

Accepted: 28 September 2025

Published: 2 October 2025

Citation: Yang, C.-C.; Wu, S.-T.; Chu, T.-W.; Liu, C.-H.; Chuang, Y.-J. Dual Machine Learning Framework for Predicting Long-Term Glycemic Change and Prediabetes Risk in Young Taiwanese Men. *Diagnostics* **2025**, *15*, 2507. <https://doi.org/10.3390/diagnostics15192507>

Copyright: © 2025 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).

Keywords: change of fasting plasma glucose; Taiwanese; machine learning; prediabetes; young men

1. Introduction

The global prevalence of type 2 diabetes (T2D) continues to rise, with an estimated 529–537 million adults affected in 2021 and projections reaching 1.31 billion by 2050 [1,2]. In Taiwan, NAHSIT 2013–2016 data show that nearly one-third of adults have impaired fasting glucose [3], yet only 4.1% of T2D patients meet the American Diabetes Association’s goals for glycated hemoglobin, blood pressure, and cholesterol [4].

The age of diabetes onset has been decreasing [5]. Because longer disease duration is strongly correlated with complications, younger patients are expected to experience more complications over their lifetime [6]. Accordingly, the American Diabetes Association recommends diabetes screening from age 35 [7].

Statistical and machine learning approaches have been used to predict diabetes or prediabetes in Chinese, Korean, and other cohorts [8–12], but most studies were cross-sectional or focused on older populations [8–10], with few targeting young adults [12]. In addition, interpretability methods such as Shapley additive explanations (SHAP) are rarely applied [13,14], and prior models typically addressed a single outcome (continuous change or binary classification). Table 1 summarizes these studies for comparison.

Table 1. Comparison of previous studies predicting prediabetes/diabetes risk versus the present study. Includes study population, design, methods, outcomes, interpretability, and key findings, highlighting how the present dual-framework approach differs.

| Study | Population | Study Design | Method(s) | Outcome | Interpretability | Key Findings | Difference from Present Study |
|------------------------|--|--------------------|--|---|------------------|---|--|
| Huang et al., 2022 [8] | Korean adults | Cross-sectional | Genetic risk score + oxidative stress score | Diabetes and prediabetes | None | Combined genetic and oxidative scores improved prediction | Not longitudinal, no ML, no SHAP |
| Wu et al., 2021 [9] | Middle-aged and elderly Chinese | Prospective cohort | Logistic regression | Prediabetes risk | None | Identified age, BMI, TG as key risk factors | Older population, single binary model |
| Yuk et al., 2022 [10] | Korean health checkup data | Cross-sectional | ML (AI algorithms) | Diabetes and prediabetes | Limited | Applied AI for screening | Cross-sectional only, no continuous δ -FPG |
| Liu et al., 2024 [12] | Young Chinese men (~5.8 y follow-up) | Prospective cohort | ML (RF, XGBoost, etc.) | Prediabetes | None | Developed predictive models | Did not use dual framework or SHAP |
| Present study (2025) | Young Taiwanese men, 18–35 y, n = 6247, mean follow-up 5.9 y | Prospective cohort | Dual ML framework (RF, SGB, XGBoost, EN vs. MLR) | Continuous δ -FPG + binary prediabetes | SHAP applied | FPG _{base} strongest predictor; BF, WBC, age, TSH, TG, LDL-C important without FPG _{base} | First dual-framework ML + SHAP interpretability in young Taiwanese men; reproducibility tested with repeated runs and CI estimates |

In contrast, our work leveraged a large, longitudinal cohort of young Taiwanese men with nearly six years of follow-up, employed a dual machine learning (ML) framework (continuous fasting plasma glucose (δ -FPG) and binary prediabetes), integrated SHAP for interpretability, and ensured reproducibility through repeated runs and confidence intervals. These strengths make our model more clinically actionable for early prediabetes risk stratification. Recent reviews further endorse SHAP and other explainable AI (XAI) tools for transparent clinical decision support [15,16].

This interpretability is particularly valuable in binary ML classification models, which predict categorical outcomes such as the presence or absence of prediabetes. Unlike continuous models that track gradual metabolic changes, binary methods align directly with diagnostic thresholds used in clinical practice [17–19], making them particularly suited

for screening programs and early intervention planning. When combined with SHAP, binary models not only deliver high predictive accuracy but also allow clear identification of the most influential factors driving a classification decision. This dual capacity—accurate prediction and transparent reasoning—makes binary ML models an essential complement to continuous prediction frameworks, ensuring both statistical robustness and clinical applicability.

In the present study, we enrolled 6247 young Taiwanese men followed up for 5.9 years. By using four different ML methods, our objectives were as follows:

1. Compare the accuracy of multiple logistic regression (MLR) and ML approaches.
2. Identify the most important risk factors for δ -FPG, defined as FPG at the end of follow-up minus FPG at baseline (FPG_{base}).
3. Use SHAP to further examine the impact and direction of each factor.

In addition to modeling δ -FPG as a continuous variable, we also aimed to develop and evaluate a categorical (binary) ML model to classify individuals as normal or prediabetes based on demographic, biochemical, and lifestyle factors, thereby aligning predictive modeling with clinically relevant diagnostic thresholds.

2. Materials and Methods

2.1. Participant and Study Design

The methods used in this study were described in our previous publication [11]. Data were obtained from the Taiwan MJ cohort, an ongoing prospective cohort based on health examinations conducted by the MJ Health Screening Centers in Taiwan [20]. These examinations include more than 100 biological indicators, such as anthropometric measurements, blood tests, and imaging studies. Each participant also completed a self-administered questionnaire covering personal and family medical history, current health status, lifestyle, physical activity, sleep habits, and dietary patterns [20].

This study was a secondary data analysis. At the time of health examinations, participants provided general consent for future anonymous research use. The database is maintained by the MJ Health Research Foundation. All or part of the data used in this study were authorized and provided by the MJ Health Research Foundation (Authorization Code: MJHRF2023007A). The interpretations and conclusions presented herein do not necessarily represent the views of the foundation. For further details, please refer to the foundation's annual technical reports [21].

The study protocol was approved by Institutional Review Board of the Kaohsiung Armed Forces General Hospital (IRB No.: KAFGHIRB 112-006, date of approval 21 June 2023). Since there was no sample collection from the patients, a short review of the IRB was approved, and no consent form was needed. A comprehensive list of variables used in the analysis is provided in Table 2, including demographic, biochemical, and lifestyle factors. All variables are reported with their original units (e.g., %, $\times 10^3/\mu\text{L}$, mg/dL). Clinically relevant transformations—such as log-transformations for triglycerides (TG), γ -glutamyl transpeptidase (γ -GT), and alkaline phosphatase (ALP), as well as derived ratios (TG/high density lipoprotein cholesterol (HDL-C), low density lipoprotein cholesterol (LDL-C)/HDL-C, waist-to-hip ratio)—were applied as part of feature engineering and are consistent with metabolic risk assessment guidelines.

Initially, 47,666 men were enrolled. For this study, we specifically selected participants aged 18–35 years, as described in the Introduction. The exclusion criteria were as follows:

1. Age < 18 and >35 years old.
2. Use of medications known to affect blood pressure, blood glucose, or blood lipids.
3. Abnormal plasma glucose level at baseline.

Table 2. Baseline demographic, biochemical, and lifestyle characteristics of the study cohort (n = 6247). Data are presented as mean \pm SE for continuous variables and n (%) for categorical variables.

| Variable | Mean \pm SE |
|--|--------------------|
| Age (year) | 27.75 \pm 5.11 |
| Years of follow-up | 5.89 \pm 4.21 |
| Body fat (%) | 21.38 \pm 5.50 |
| White blood cell count ($\times 10^3/\mu\text{L}$) | 6.23 \pm 1.44 |
| Hemoglobin ($\times 10^6/\mu\text{L}$) | 15.43 \pm 0.99 |
| Platelets ($\times 10^3/\mu\text{L}$) | 236.72 \pm 49.57 |
| Fasting plasma glucose-baseline (mg/dL) | 92.01 \pm 4.74 |
| Glutamic pyruvic transaminase (IU/L) | 31.51 \pm 47.78 |
| Glutamic oxaloacetic transaminase (IU/L) | 24.13 \pm 20.85 |
| γ -Glutamyl transpeptidase (IU/L) | 19.88 \pm 16.96 |
| Lactate dehydrogenase (IU/L) | 287.87 \pm 66.74 |
| Uric acid (mg/dL) | 7.08 \pm 1.40 |
| Creatinine (mg/dL) | 1.08 \pm 0.13 |
| Triglyceride (mg/dL) | 100.36 \pm 60.99 |
| High-density lipoprotein cholesterol (mg/dL) | 49.25 \pm 11.87 |
| Low-density lipoprotein cholesterol (mg/dL) | 112.56 \pm 31.15 |
| Alkaline phosphatase (IU/L) | 147.33 \pm 47.38 |
| Thyroid stimulating hormone (IU/mL) | 1.61 \pm 1.66 |
| C-reactive protein (mg/dL) | 0.21 \pm 0.40 |
| Drinking area | 1.60 \pm 7.29 |
| Sport area | 9.55 \pm 9.03 |
| δ -Fasting plasma glucose (mg/dL) | 5.21 \pm 7.02 |
| Spouse status | |
| Single | 3957 (63.9%) *** |
| With spouse | 2232 (36.1%) |
| Sleep hours | |
| 0–4 h/day | 24 (0.4%) |
| 4–6 h/day | 1054 (16.9%) |
| 6–8 h/day | 4745 (76.2%) |
| >8 h/day | 408 (6.5%) |
| Education level | |
| Primary school | 3 (0.1%) |
| Junior high school | 51 (0.8%) |
| Senior high school | 1012 (16.2%) |
| College | 1830 (29.3%) |
| University | 2293 (36.8%) |
| Higher than master's degree | 1031 (16.8%) |
| Income level (thousand USD/year) | |
| 0 | 1232 (19.7%) * |
| 12.7/year | 1029 (16.5%) |
| 12.7–25.3/year | 2822 (45.2%) |
| 25.3–38.0/year | 883 (14.1%) |
| 38.0–50.6/year | 130 (2.1%) |
| 50.6–63.3/year | 73 (1.2%) |
| >63.3/year | 78 (1.2%) |

p: * < 0.05, *** p: < 0.001.

After excluding subjects who did not fit our inclusion criteria, 6247 men were left for further analysis (Figure 1). Methods of how to collect demographic, biochemistry, and lifestyle data were published by our group previously and are not shown here [19].




Figure 1. The scheme of participant selection.

Table 2 shows the 25 baseline variables. These included participants' age, body fat (BF), complete blood cell count, biochemistries, thyroid stimulating hormone (TSH), c-reactive protein (CRP), education level, relationship status, and income level. The drinking area was defined as the multiple of total drinking duration, frequency of drinking, and the percentage of alcohol. Similarly, the smoking area was the multiple of the duration, frequency of smoking, and number of cigarettes. The sport area was the multiple of duration, frequency, and type of exercise. All the aforementioned parameters were the independent variables, and the dependent variable was δ -FPG after an average of 5.9 years follow-up.

Two models were built in the present study. From our preliminary evaluation, Model 1 included all 25 variables. Our results showed that FPG_{base} had 100% importance compared to the second important factor (BF, 17.6%). In order to further evaluate the hidden interactions between these factors, Model 2 was built by removing the FPG_{base} .

2.2. Traditional Statistical Method

MLR was used, with δ -FPG as the dependent variable and demographic, biochemical, and lifestyle variables as independent predictors. Further technical details were provided in our previous publication [12].

2.3. Machine Learning Methods

We employed two complementary ML frameworks to predict prediabetes risk and identify key predictors from demographic, biochemical, and lifestyle data. The first modeled the outcome as a continuous variable— δ -FPG—while the second modeled a binary outcome (normal vs. prediabetes). This dual design allowed us to capture both fine-grained metabolic variation and clinically actionable thresholds.

2.3.1. Continuous Outcome Prediction (δ -FPG)

Four ML algorithms—random forest (RF) [22], stochastic gradient boosting (SGB) [23], eXtreme gradient boosting (XGBoost) [24], and elastic net (EN) [25]—were trained to predict δ -FPG (Figure 2). Data were split 80/20 into training and testing sets. Hyperparameters were tuned via 10-fold cross-validation with grid search. Performance was assessed with symmetric mean absolute percentage error (SMAPE), relative absolute error (RAE), root relative squared error (RRSE), and root mean squared error (RMSE) (Table 3).

To ensure robustness, the training/testing procedure was repeated 10 times with different random seeds. The mean and 95% confidence intervals (CIs) of performance metrics across these runs demonstrate model stability and reproducibility. Feature importance scores were averaged across all runs and models. Feature importance distributions were examined using boxplots of SHAP-derived importance, generated from 20 repeated training–testing cycles with different random seeds. Narrow boxplots indicate stable contributions across runs, whereas wider distributions highlight features with unstable or

context-dependent importance. This procedure has been recommended for evaluating the reliability of explanation methods in machine learning [26].




Figure 2. Proposed machine learning prediction scheme.

Table 3. Equations and descriptions of performance metrics used to evaluate continuous outcome prediction models: SMAPE, RAE, RRSE, and RMSE.

| Metrics | Description | Calculation |
|---------|--|---|
| SMAPE | Symmetric mean absolute percentage error | $SMAPE = \frac{1}{n} \sum_{i=1}^n \frac{ y_i - \hat{y}_i }{(y_i + \hat{y}_i)/2} \times 100$ |
| RAE | Relative absolute error | $RAE = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y}_i)^2}}$ |
| RRSE | Root relative squared error | $RRSE = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y}_i)^2}}$ |
| RMSE | Root mean squared error | $RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$ |

where \hat{y}_i and y_i represent predicted and actual values, respectively; n stands the number of instances.

A planned sensitivity analysis removed FPG_{base} due to its dominant influence (100% importance in initial models). The same modeling procedures were applied to Model 2 (without FPG_{base}) to identify alternative predictors. Feature importance for both Model 1 and Model 2 are presented in Section 3.3.

Model interpretability for the FPG_{base}-excluded model (Model 2) was further explored using SHAP applied to the random forest algorithm. The SHAP bee swarm plot visualizes the direction and magnitude of each feature's impact on individual predictions. The mean absolute SHAP values provide a global ranking of feature importance, while the net SHAP values indicate whether features exert predominantly positive or negative effects overall. Dependence plots for key features (e.g., follow-up interval and baseline FPG) are also shown in Section 3.3, illustrating how feature effects are modulated by other variables.

Residual diagnostics for the continuous models, including quantile–quantile plots and residuals vs. fitted values, were used to validate model assumptions.

2.3.2. Binary Outcome Prediction (Normal vs. Prediabetes)

For classification, we developed an XGBoost-based pipeline [24] using a 70/30 derivation/validation split. Feature engineering included clinically relevant ratios (TG/HDL-C, LDL-C/HDL-C, waist-to-hip ratio) and log transformations (TG, γ-GT, ALP). Missing data were imputed using median values [27]. Class imbalance was addressed with SMOTE-Tomek [28,29]. We also considered recent evidence and reviews on imbalance remedies to mitigate over-optimism and miscalibration risks [30,31]. Feature selection used permutation importance from a baseline XGBoost model, retaining the top 20 predictors. Hyperparameters were tuned with Optuna [32] in threefold stratified CV. Final model performance was evaluated on the validation set using accuracy, precision, recall, specificity, F1 score, ROC-AUC, and PR-AUC. Model interpretability was assessed with SHAP [13]. The confusion matrix at a threshold of 0.50 is also shown. Given class imbalance, precision-recall curves were prioritized for conveying retrieval performance [33]. All preprocessing steps—including median imputation, SMOTE-Tomek resampling, feature engineering, and feature selection—were applied independently within each training fold during cross-validation to prevent data leakage. Hyperparameter tuning was performed using Optuna with threefold stratified CV, 100 trials, early stopping after 20 rounds, and a fixed random seed (42) for reproducibility.

To prevent data leakage, all preprocessing steps—including missing value imputation, feature engineering, SMOTE-Tomek resampling, and feature selection—were applied independently within each training fold during cross-validation. Specifically, the median for imputation was computed solely from the training fold and applied to both training and validation subsets. Similarly, SMOTE-Tomek was applied only to the training fold, without incorporating any validation data. Feature selection based on permutation importance was performed using a model trained on the training fold only. Model calibration was assessed using the Brier score and a reliability (calibration) curve, where predicted probabilities were compared with observed outcomes across 10 bins of predicted risk, providing a visual assessment of calibration. Brier score is defined as

$$\text{Brier Score} = \frac{1}{n} \sum_{i=1}^n (\hat{p}_i - y_i)^2$$

where \hat{p}_i is the predicted probability of prediabetes for the i -th individual, y_i is the observed outcome (1 if prediabetic, 0 otherwise), and n is the number of individuals in the validation set. A lower Brier score indicates better calibration, with 0 representing perfect calibration. These choices align with recent BMJ guidance on evaluating clinical prediction models, covering discrimination, calibration, and external validation [34,35].

To evaluate the clinical utility of the model across different decision contexts, we assessed two operating thresholds: (1) a high-sensitivity threshold (0.2892) optimized to minimize false negatives, which is appropriate for screening settings, and (2) a balanced threshold (0.5683) identified using Youden's J statistic, which maximizes the sum of sensitivity and specificity. The optimal threshold was determined by maximizing Youden's J = (sensitivity + specificity – 1). Decision-curve analysis (DCA) was performed to quantify the net benefit of the model across a range of threshold probabilities, comparing it against strategies of treating all or none. DCA was performed to evaluate the clinical utility of the binary classifier by quantifying the net benefit of using the model to guide clinical decisions across a range of threshold probabilities [36,37]. The net benefit was calculated as the difference between the proportion of true positives and the weighted proportion

of false positives, where the weight is the threshold probability divided by (1-threshold probability). This allowed the comparison of the model's performance against two extreme strategies: treating all patients or treating none. DCA was implemented using the predicted probabilities from the validation set and the true outcome labels, across threshold probabilities ranging from 0.01 to 0.99.

2.3.3. Rationale for Combined Framework

The continuous framework offers precise measurement of glycemic change, while the binary framework aligns with diagnostic thresholds for prediabetes. Presenting both ensures statistical rigor and clinical relevance.

2.3.4. Confidence Interval Estimation

To quantify the uncertainty around model performance, we calculated 95% CIs for each evaluation metric (SMAPE, RMSE, RAE, RRSE) based on the distribution of scores obtained from 10 repeated training–testing cycles for each model. For each metric, the mean and standard error (SE) were computed, and the CIs was estimated as:

$$CIs = \bar{x} \pm t\alpha/2, n - 1 \times SE$$

CIs = $\bar{x} \pm t\alpha/2, n - 1 \times SE$, where \bar{x} is the sample mean of the metric, $t\alpha/2$; $n - 1$ is the critical value from the Student's t-distribution with $n - 1$ degrees of freedom [38]; and SE is the standard error of the mean. This approach follows established statistical practice for small-sample inference [39] and is applicable when the sampling distribution of the metric is approximately normal [40].

2.3.5. Sensitivity and Ablation Analysis

To assess model robustness and the concentration of predictive signal, we performed an ablation analysis by removing the top 3 most important features from the binary XGBoost model. The resulting drop in ROC-AUC is reported in Supplementary Table S1, illustrating the model's dependence on key predictors.

2.3.6. Multicollinearity Assessment

To ensure feature independence and model stability, we computed the variance inflation factor (VIF) for all predictors. Results are presented in Supplementary Table S2, with all VIF values below 10, indicating acceptable levels of multicollinearity. The feature correlation heatmap is shown in Supplementary Figure S1. These diagnostic tools help ensure that observed feature importance is not an artifact of highly correlated variables, which can distort model interpretability in biomedical datasets [41].

2.3.7. Reproducibility Artifact

To enhance transparency and reproducibility, we provided a minimal, self-contained Python v3.10.12 script that generated synthetic data mimicking the structure of our cohort ($n = 100$ young Taiwanese men with features including age, body fat, WBC, TSH, TG, LDL-C, and δ -FPG) and executed a simplified version of our modeling pipeline (Supplementary Figure S2). The script trained an XGBoost Regressor to predict δ -FPG, and it computed SHAP values for interpretability and RMSE as a performance metric. The synthetic data were generated using biologically plausible distributions and correlations, ensuring the code demonstrated the core workflow without requiring access to the real, privacy-sensitive MJ Health data. This artifact is available from the corresponding author upon reasonable request and fulfills the journal's reproducibility guidelines.

3. Results

3.1. Demographic Characteristics

Our analysis included 6247 young Taiwanese men aged 18–35 years at baseline, with a mean follow-up duration of 5.9 years (Figure 1). The selection scheme and exclusion criteria are detailed in Figure 1. Baseline demographic, biochemical, and lifestyle characteristics of the full cohort are summarized in Table 2.

To contextualize our work, Table 1 compares our study design, population, methods, and findings with four recent studies predicting prediabetes or diabetes risk. Our study is distinguished by its focus on young men, prospective design, dual ML framework (continuous δ -FPG + binary classification), application of SHAP interpretability, and rigorous assessment of reproducibility through repeated runs and CIs.

Participants were stratified by glycemic outcome at follow-up: Group 1 ($n = 2789$) developed prediabetes, while Group 2 ($n = 3458$) remained normoglycemic. Significant differences between groups are shown in Table 4. Compared to Group 2, Group 1 participants were older (28.4 vs. 27.3 years, $p < 0.001$) and had higher body fat (22.1 vs. 20.8 mg/dL, $p < 0.001$), higher white blood cell count (6.3 vs. $6.2 \times 10^3/\mu\text{L}$, $p < 0.001$), higher baseline FPG (93.3 vs. 91.0 mg/dL, $p < 0.001$), higher triglycerides (105.8 vs. 96.0 mg/dL, $p < 0.001$), and higher LDL-C (115.0 vs. 110.6 mg/dL, $p < 0.001$). They also had shorter follow-up duration (5.6 vs. 6.1 years, $p < 0.001$) and lower HDL-C (48.4 vs. 50.0 mg/dL, $p < 0.001$). Marital status (fewer single men in prediabetes group) and income level (fewer unemployed in prediabetes group) also differed significantly.

Table 4. Demographic and biochemical characteristics of participants stratified by glycemic outcome (prediabetes and normal). Continuous variables are shown as mean \pm SE and categorical variables as n (%).

| Variable | Prediabetes | Normal |
|--|------------------|----------------------|
| n | 2789 | 3458 |
| Age (year) | 28.4 ± 5.0 | 27.3 ± 5.1 *** |
| Years of follow-up | 5.6 ± 4.0 | 6.1 ± 4.4 *** |
| Body fat (%) | 22.1 ± 5.5 | 20.8 ± 5.5 *** |
| White blood cell count ($\times 10^3/\mu\text{L}$) | 6.3 ± 1.5 | 6.2 ± 1.4 *** |
| Hemoglobin ($\times 10^6/\mu\text{L}$) | 15.5 ± 1.0 | 15.4 ± 1 |
| Platelets ($\times 10^3/\mu\text{L}$) | 237.1 ± 49.8 | 236.4 ± 49.4 |
| Fasting plasma glucose—baseline (mg/dL) | 93.3 ± 4.3 | 91.0 ± 4.8 *** |
| Glutamic pyruvic transaminase (IU/L) | 32.7 ± 50.6 | 30.5 ± 45.4 |
| Glutamic oxaloacetic transaminase (IU/L) | 24.2 ± 20.2 | 24.1 ± 21.4 |
| γ -Glutamyl transpeptidase (IU/L) | 20.6 ± 17.1 | 19.3 ± 16.8 *** |
| Lactate dehydrogenase (IU/L) | 290.6 ± 63 | 285.6 ± 69.6 *** |
| Uric acid (mg/dL) | 7.2 ± 1.4 | 7.0 ± 1.4 *** |
| Creatinine (mg/dL) | 1.1 ± 0.1 | 1.1 ± 0.1 |
| Triglyceride (mg/dL) | 105.8 ± 66.1 | 96.0 ± 56.2 *** |
| High-density lipoprotein cholesterol (mg/dL) | 48.4 ± 11.6 | 50.0 ± 12.1 *** |
| Low-density lipoprotein cholesterol (mg/dL) | 115.0 ± 32.0 | 110.6 ± 30.3 *** |
| Alkaline phosphatase (IU/L) | 146.7 ± 45.0 | 147.8 ± 49.3 |
| Thyroid-stimulating hormone (IU/mL) | 1.6 ± 0.8 | 1.7 ± 2.1 * |
| C-reactive protein (mg/dL) | 0.2 ± 0.4 | 0.2 ± 0.4 |
| Drinking area | 1.8 ± 7.4 | 1.5 ± 7.2 |
| Sport area | 9.2 ± 8.9 | 9.8 ± 9.1 * |
| δ -Fasting plasma glucose (mg/dL) | 9.88 ± 5.45 | 1.44 ± 5.77 *** |
| Spouse status | | |
| Single | 1660 (60.1%) | 2297 (67.1%) *** |
| With spouse | 1104 (39.9%) | 1128 (32.9%) |

Table 4. Cont.

| Variable | Prediabetes | Normal |
|----------------------------------|--------------|-----------------|
| Sleep hours | | |
| 0–4 h/day | 8 (0.3%) | 16 (0.5%) |
| 4–6 h/day | 474 (17.0%) | 580 (16.8%) |
| 6–8 h/day | 2129 (76.6%) | 2616 (75.8%) |
| >8 h/day | 170 (6.1%) | 238 (6.9%) |
| Education level | | |
| Primary school | 1 (0.04%) | 2 (0.1%) * |
| Junior high school | 18 (0.7%) | 33 (1.0%) |
| Senior high school | 425 (15.3%) | 587 (17.1%) |
| College | 835 (30.0%) | 995 (28.9%) |
| University | 1006 (36.2%) | 1287 (37.4%) |
| Higher than master's degree | 496 (17.8%) | 535 (15.6%) |
| Income level (thousand USD/year) | | |
| 0 | 471 (16.9%) | 761 (22.0%) *** |
| 12.7/year | 430 (15.4%) | 599 (17.3%) |
| 12.7–25.3/year | 1292 (46.3%) | 1530 (44.3%) |
| 25.3–38.0/year | 448 (16.1%) | 435 (12.6%) |
| 38.0–50.6/year | 70 (2.5%) | 60 (1.7%) |
| 50.6–63.3/year | 42 (1.5%) | 31 (0.9%) |
| >63.3/year | 36 (1.3%) | 42 (1.2%) |

Note: Prediabetes: subjects became prediabetic at the end of follow-up; Normal: subjects remained a normal glycemic level at the end of follow-up. p : * < 0.05 , *** p : < 0.001 .

3.2. Continuous Outcome Prediction (δ -FPG)

We first evaluated the ability of four machine learning models: RF, SGB, XGBoost, and EN, to predict the continuous outcome δ -FPG. Performance metrics (SMAPE, RAE, RRSE, RMSE) and their definitions are provided in Table 5.

Table 5. Average performance of linear regression and machine learning methods for predicting continuous glycemic change (δ -FPG), with and without FPG_{base}.

| With FPG _{base} | | | | |
|-----------------------------|--------|--------|--------|--------|
| Methods | SMAPE | RAE | RRSE | RMSE |
| MLR | 0.968 | 0.9283 | 0.9077 | 6.4156 |
| RF | 0.9571 | 0.9273 | 0.9074 | 6.4133 |
| SGB | 0.9426 | 0.9306 | 0.9113 | 6.4407 |
| XGBoost | 0.9443 | 0.929 | 0.9102 | 6.4329 |
| EN | 0.9599 | 0.9266 | 0.9068 | 6.4092 |
| Without FPG _{base} | | | | |
| Methods | SMAPE | RAE | RRSE | RMSE |
| MLR | 0.951 | 0.9845 | 0.9832 | 6.9022 |
| RF | 0.9477 | 0.9887 | 0.9854 | 6.9175 |
| SGB | 0.9489 | 0.9839 | 0.9817 | 6.8916 |
| XGBoost | 0.9498 | 0.9878 | 0.985 | 6.9149 |
| EN | 0.9492 | 0.9845 | 0.9827 | 6.8985 |

Note: MLR: multiple linear regression; RF: random forest; SGB: stochastic gradient boosting; XGBoost: eXtreme Gradient Boosting; EN: elastic net; SMAPE: symmetric mean absolute percentage error; RAE: relative absolute error; RRSE: root relative squared error; RMSE: root mean squared error.

Model 1 (all predictors included): All ML models outperformed MLR. Elastic net achieved the lowest RMSE (6.4092) and RRSE (0.9068), followed closely by random forest (6.4133, 0.9074) and XGBoost (6.4329, 0.9102) (Table 5).

Model 2 (FPG_{base} excluded): As expected, performance declined slightly across all models due to removal of the strongest predictor. EN performed best (RMSE: 6.8985, RRSE: 0.9827), followed by SGB (RMSE: 6.8916, RRSE: 0.9817) and RF (RMSE: 6.9175, RRSE: 0.9854) (Table 5).

To assess reproducibility, we repeated the training/testing procedure 10 times. Table 6 presents the mean and 95% confidence intervals for each metric. The intervals were extremely narrow (often identical upper and lower bounds), indicating high stability and reproducibility of model performance across runs. For example, RF RMSE was consistently mean = 6.471158, SE = 0.00179, 95% CI = 6.467109–6.475208.

Table 6. Mean and 95% confidence intervals (CIs) of performance metrics (SMAPE, RMSE, RAE, RRSE) across 10 repeated runs for continuous outcome prediction models, demonstrating reproducibility and stability.

| Model | Metric | Mean | SE | 95% CI Lower | 95% CI Upper |
|---------|--------|----------|----------|--------------|--------------|
| EN | SMAPE | 96.46007 | 0.016408 | 96.42295 | 96.49719 |
| EN | RMSE | 6.427868 | 0.00196 | 6.423434 | 6.432301 |
| EN | RAE | 0.922647 | 0.000325 | 0.921912 | 0.923382 |
| EN | RRSE | 0.905753 | 0.00021 | 0.905279 | 0.906227 |
| RF | SMAPE | 96.84203 | 0.021496 | 96.7934 | 96.89066 |
| RF | RMSE | 6.471158 | 0.00179 | 6.467109 | 6.475208 |
| RF | RAE | 0.920818 | 0.000309 | 0.920119 | 0.921517 |
| RF | RRSE | 0.912466 | 0.000341 | 0.911695 | 0.913237 |
| XGBoost | SMAPE | 104.8296 | 0.024476 | 104.7742 | 104.8849 |
| XGBoost | RMSE | 7.056346 | 0.002225 | 7.051313 | 7.06138 |
| XGBoost | RAE | 1.004276 | 0.000322 | 1.003548 | 1.005003 |
| XGBoost | RRSE | 0.994844 | 0.000363 | 0.994023 | 0.995666 |

Note: SMAPE: symmetric mean absolute percentage error; RMSE: root mean squared error; RAE: relative absolute error; RRSE: root relative squared error. RF: random forest; XGBoost: eXtreme Gradient Boosting, EN: elastic net.

Residual diagnostics for the continuous models, including quantile–quantile plots and residuals vs. fitted values, are presented in Figure 3, confirming reasonable adherence to model assumptions.




Figure 3. (A) Residuals vs. fitted values for Linear Regression (red). (B) Residuals vs. fitted values for XGBoost Regression (cyan). The horizontal dashed line at zero indicates no residual bias. Patterns or funnel shapes may suggest heteroscedasticity or model misspecification.

3.3. Feature Importance and Model Interpretability (Continuous Models)

In the full model (Model 1), which included all baseline predictors, FPG_{base} had the maximum possible importance score of 100% across all four ML methods (Table 7, Figure 4). This overwhelming dominance of FPG_{base} meant that its predictive contribution far exceeded all other variables, with the second-ranked feature, BF, having a mean importance of only 17.6%. Such a strong single-variable effect can mask the influence of other clinically relevant predictors.

Table 7. Relative importance percentages of predictors in Model 1 (with FPG_{base} included) across four machine learning methods (RF, SGB, XGBoost, EN).

| Variables | MOIP |
|---|-------|
| Baseline fasting plasma glucose | 100 |
| Body fat | 17.64 |
| Thyroid-stimulating hormone | 14.8 |
| White blood cell count | 14.51 |
| Years of follow-up | 14.32 |
| C-reactive protein | 11.63 |
| Age | 10.30 |
| HDL-cholesterol (HDL-C) | 9.75 |
| Uric acid | 9.34 |
| Hemoglobin | 8.80 |
| Lactic dehydrogenase | 8.32 |
| LDL-cholesterol (LDL-C) | 8.30 |
| Triglyceride (TG) | 7.63 |
| Alkaline phosphatase | 7.39 |
| Platelet | 7.11 |
| Glutamic pyruvic transaminase (GPT) | 5.93 |
| γ-Glutamyl transpeptidase | 5.27 |
| Spouse status | 4.87 |
| Glutamic oxaloacetic transaminase (GOT) | 4.50 |
| Sport area | 4.27 |
| Income status | 3.50 |
| Creatinine | 2.41 |
| Sleeping level | 2.15 |
| Education level | 1.71 |
| Drinking area | 0.54 |

Note: RF: random forest; SGB: stochastic gradient boosting; XGBoost: eXtreme Gradient Boosting; EN: elastic net; MOIP: mean of importance percentage across four ML models (RF, SGB, XGBoost, EN). Body fat: adiposity marker—consistently pro-glycemic. White blood cell count: inflammatory marker—higher levels indicate chronic low-grade inflammation linked to insulin resistance. Thyroid-stimulating hormone: thyroid function marker—higher levels associated with reduced metabolic rate and increased adiposity. TG, LDL-C, HDL-C: lipid panel—core components of insulin resistance dyslipidemia (high TG, high LDL-C, low HDL-C). Age: aging marker—associated with declining β-cell function and increasing insulin resistance.

To address this, we developed a second model (Model 2) with FPG_{base} removed from the predictor set. In this model, BF emerged as the top predictor (mean importance: 88.85%), followed by white blood cell count (WBC, 58.62%), age (54.89%), thyroid-stimulating hormone (TSH, 36.87%), triglycerides (TG, 32.66%), and LDL-C (28.69%) (Table 8, Figure 5). The redistribution of importance scores revealed a broader set of influential features that may otherwise be overlooked when FPG_{base} is included.




Figure 4. The graphic demonstration of the importance percentage in Model 1 (with FPG_{base} included). The darker green bars show the most important six features of the study.

Table 8. Relative importance percentages of predictors in Model 2 (excluding FPG_{base}) across four machine learning methods (RF, SGB, XGBoost, EN).

| Variables | MOIP |
|---|-------|
| Body fat | 88.85 |
| White blood cell count | 58.62 |
| Age | 54.89 |
| Thyroid stimulating hormone | 36.87 |
| Triglyceride | 32.66 |
| LDL-cholesterol | 28.69 |
| HDL-cholesterol | 27.42 |
| Years of follow-up | 24.94 |
| γ -Glutamyl transpeptidase | 22.78 |
| Hemoglobin | 22.67 |
| Lactic dehydrogenase | 22.32 |
| Alkaline phosphatase | 22.16 |
| Platelet | 21.36 |
| Uric acid | 20.95 |
| Glutamic oxaloacetic transaminase (GOT) | 18.88 |
| Glutamic pyruvic transaminase (GPT) | 18.49 |
| Sport area | 16.83 |
| Creatinine | 8.95 |
| Income level | 6.71 |
| Education level | 6.11 |
| C-reactive protein | 5.14 |
| Sleeping level | 4.00 |
| Drink area | 1.65 |
| Spouse status | 0 |

Note: RF: random forest; SGB: stochastic gradient boosting; XGBoost: eXtreme Gradient Boosting; EN: elastic net; MOIP: mean of importance percentage across four ML models (RF, SGB, XGBoost, EN). Body fat: adiposity marker—consistently pro-glycemic. White blood cell count: inflammatory marker—higher levels indicate chronic low-grade inflammation linked to insulin resistance. Thyroid-stimulating hormone: thyroid function marker—higher levels associated with reduced metabolic rate and increased adiposity. TG, LDL-C, HDL-C: lipid panel—core components of insulin resistance dyslipidemia (high TG, high LDL-C, low HDL-C). Age: aging marker—associated with declining β -cell function and increasing insulin resistance.




Figure 5. The graphic demonstration of the importance percentage in Model 2 (excluding FPG_{base}). The darker green bars show the most important six features of the study.

Although the exclusion of FPG_{base} led to a slight decline in predictive performance across all ML methods, the results remained stable, suggesting that these secondary predictors retained substantial explanatory value for $\delta\text{-FPG}$. This sensitivity analysis highlights that while FPG_{base} is the single strongest predictor, other metabolic, inflammatory, and lipid-related factors also play significant independent roles in determining future changes in fasting plasma glucose.

Model interpretability was assessed using SHAP applied to the RF model for the FPG_{base} -excluded dataset. Figure 6 presents the bee swarm plot, showing the distribution and direction of SHAP values for each predictor. Variables such as years of follow-up, TSH, HDL-C, age, LDH, and BF had the most substantial influence on $\delta\text{-FPG}$, with color gradients indicating whether higher feature values were associated with increases (red) or decreases (blue) in predicted $\delta\text{-FPG}$.

Figure 7 shows the mean absolute SHAP values for each feature, ranking them by overall importance. Years of follow-up, TSH, HDL-C, and age remained the top contributors, confirming their consistent predictive role.

Figure 8 displays the net SHAP values, reflecting the difference between the mean positive and negative contributions for each feature. This plot highlights features whose effects are unidirectional (consistently increasing or decreasing $\delta\text{-FPG}$) versus those with bidirectional influence depending on their value. Together, these three visualizations provide a comprehensive understanding of both the magnitude and directionality of each predictor's impact on model predictions.

There are two key SHAP dependence plots that illustrate feature interactions. Figure 9 shows the effect of follow-up interval (Yr-GAP) on SHAP values was modulated by FPG_{base} . Longer follow-up generally increased risk, but this effect was stronger in individuals with lower FPG_{base} . Figure 10 shows the effect of FPG_{base} was modulated by body fat percentage. Lower FPG_{base} was associated with higher risk (positive SHAP), especially in individuals with higher body fat.




Figure 6. SHAP bee swarm plot of the random forest model, illustrating the distribution and magnitude of predictor effects (red = higher values, blue = lower values). Note: Red dots (higher feature value) push predictions upward (worsening glycemia); blue dots (lower feature value) push predictions downward (improving glycemia). For example, higher body fat consistently increases predicted glucose rise, while higher HDL-C is protective.




Figure 7. Absolute SHAP value ranking of predictors, showing the magnitude of each variable's contribution to glycemic trajectory.




Figure 8. Net SHAP value plot of predictors, demonstrating the overall directional impact of features on glycemic outcomes.




Figure 9. SHAP dependence plot for follow-up interval (Yr-GAP), colored by FPG_{base} .




Figure 10. SHAP dependence plot for FPG_{base} , colored by body fat percentage. Each point represents an individual participant. The x -axis shows FPG_{base} , and the y -axis indicates its SHAP value. Lower FPG_{base} levels were associated with strong positive SHAP values, while higher levels showed negative contributions. Body fat percentage, represented by the color gradient, further modulated this relationship.

Each point represents an individual participant. The x -axis shows years between baseline and follow-up examinations (Yr-GAP), and the y -axis indicates the corresponding SHAP value. A longer follow-up interval was associated with progressively higher SHAP values, suggesting an increasing contribution to the model prediction. The color gradient denotes FPG_{base} , which modulated the effect size.

3.4. Binary Outcome Prediction (Prediabetes vs. Normal)

The binary classification model was evaluated for its ability to distinguish between normal and prediabetes cases. Overall model performance is illustrated in the precision-recall curve (Figure 11) and the receiver operating characteristic curve (Figure 12), with corresponding metrics summarized in Table 9. The categorical model, developed using the top 20 predictors identified by permutation importance, demonstrated an overall accuracy of 0.788 and a precision of 0.791 (Table 9). Sensitivity was exceptionally high (0.995), correctly identifying almost all prediabetes cases, whereas specificity was low (0.010), indicating a tendency to classify most individuals as prediabetes. The F1 score was 0.881, reflecting a strong balance between precision and recall in this high-sensitivity context. The model achieved a PR-AUC of 0.873, highlighting good performance in ranking positive cases, while the ROC-AUC was more modest at 0.667, consistent with the class imbalance and prioritization of recall. Out of 1477 actual prediabetes cases, the model correctly classified 1473 (true positives) and missed only 4 (false negatives), but also incorrectly labeled 390 normal individuals as prediabetes (false positives). Supplementary Table S1 reports 95% bootstrap confidence intervals for ROC-AUC (0.6697–0.6882) and PR-AUC (0.6349–0.6630), quantifying the uncertainty around these estimates.




Figure 11. Precision–recall curve. Note: Precision–recall curve of the binary classifier. High recall confirmed excellent sensitivity for prediabetes detection.




Figure 12. The plot of ROC curve. Note: ROC curve of the binary classifier. Demonstrates moderate discrimination, consistent with prioritization of sensitivity. The black dotted line represents the performance of a random classifier ($AUC = 0.5$) which indicates no significance.

Table 9. Performance metrics of the categorical (binary) model for predicting prediabetes versus normal outcomes, including accuracy, precision, recall, specificity, false positive rate, F1 score, ROC-AUC, PR-AUC, and classification counts.

| Accuracy | Precision | Sensitivity (Recall) | Specificity | False Positive Rate | F1 Score | ROC- AUC | PR-AUC | True Positives | False Positives | True Negatives | False Negatives |
|----------|-----------|-------------------------|-------------|---------------------------|-------------|-------------|--------|-------------------|--------------------|-------------------|--------------------|
| 0.788 | 0.791 | 0.995 | 0.010 | 0.990 | 0.881 | 0.667 | 0.873 | 1473 | 390 | 4 | 8 |

At a probability threshold of 0.50, the confusion matrix (Figure 13) shows 1473 true positives, 390 false positives, 4 true negatives, and 8 false negatives.




Figure 13. Confusion matrix at chosen threshold. Confusion matrix at the prespecified threshold of 0.50. Cell values indicate the number of true positives, false positives, true negatives, and false negatives.

The calibration plot (Figure 14) showed predicted risks aligned well with observed outcomes across bins. It demonstrated that the model's predicted probabilities closely aligned with observed outcomes, particularly at higher risk levels (>0.6), though slight under-prediction was observed at lower probabilities.




Figure 14. Calibration plot (reliability curve) of the XGBoost model for predicting prediabetes.

3.5. Clinical Utility and Threshold Analysis

To support clinical utility in screening settings, we evaluated a high-sensitivity operating threshold that prioritizes minimizing false negatives. At a threshold of 0.2892, the model achieved a sensitivity of 99.53% and specificity of 47.46% (Table 10). This high-sensitivity point is suitable for population-level screening, where identifying all at-risk individuals is critical, even at the cost of increased false positives.

Table 10. Performance of the binary model at a high-sensitivity threshold (0.2892), prioritizing minimization of false negatives for screening purposes.

| Metric | Value |
|-------------|--------|
| Threshold | 0.2892 |
| Sensitivity | 0.9953 |
| Specificity | 0.4746 |

For diagnostic settings requiring a balance between sensitivity and specificity, we identified an optimal balanced threshold using Youden's J statistic. The threshold of 0.5683 maximized Youden's J (0.7930), yielding a sensitivity of 88.69% and specificity of 90.61% (Table 11). This point provides a favorable trade-off between detection and precision, making it appropriate for confirmatory assessment or targeted intervention programs.

Table 11. Performance of the binary model at the optimal balanced threshold (0.5683), determined by Youden's J statistic, representing a trade-off between sensitivity and specificity suitable for diagnostic settings.

| Metric | Value |
|-------------|--------|
| Threshold | 0.5683 |
| Sensitivity | 0.8869 |
| Specificity | 0.9061 |
| Youden's J | 0.7930 |

To assess the clinical utility of our prediction model, we performed a DCA. As shown in Figure 15, the RF model provides a greater net benefit than both the 'treat all' and 'treat none' strategies across a wide range of threshold probabilities (10% to 80%). This indicates that using the model to guide screening and intervention decisions would lead to better patient outcomes than either universal screening or no screening. The model's clinical utility is highest in the moderate risk range, supporting its use as a tool for targeted prevention in young men.




Figure 15. Decision curve analysis. Note: The DCA model provided superior net clinical benefit compared to 'treat all' and 'treat none'.

To assess probabilistic calibration, we computed the Brier score, which was 0.1754. This indicates good to moderate calibration, particularly given the inherent uncertainty in predicting long-term prediabetes risk in a relatively healthy young population. We evaluated the performance of the XGBoost model at two clinically relevant operating points. At a high-sensitivity threshold of 0.2892, the model achieved a sensitivity of 99.53% and specificity of 47.46%, minimizing false-negative classifications and making it suitable for population-level screening. In contrast, the balanced threshold (determined by Youden's $J = 0.7930$) was set at 0.5683, yielding a sensitivity of 88.69% and specificity of 90.61%, representing an optimal trade-off between sensitivity and specificity for diagnostic use.

3.6. Model Robustness and Sensitivity Analysis

To assess the concentration of predictive signal, we removed the top three most important features from the binary XGBoost model. This reduced the ROC-AUC from 0.6753 to 0.5588 (Table 12), indicating that while the model relies on key predictors, it retains some predictive capacity even without them.

Table 12. Sensitivity and ablation analysis.

| Model | ROC-AUC |
|----------------------------------|---------|
| Original XGBoost | 0.6753 |
| XGBoost (top 3 features removed) | 0.5588 |

VIF values for all predictors were below 10 (max VIF = 8.00 for x8), indicating acceptable levels of multicollinearity (Supplementary Table S2). The feature correlation heatmap is shown in Supplementary Figure S2.

4. Discussion

In the present study, we applied four machine learning (ML) methods to identify the most important risk factors for developing prediabetes after an average follow-up of 5.9 years. FPG_{base} emerged as the dominant predictor, with an importance score of 100% in Model 1. The second most important factor, BF, reached only 17.6%. Because FPG_{base} could overshadow other variables, we constructed Model 2 by excluding it. In this model, BF, WBC, age, TSH, LDL-C, and HDL-C became the leading determinants. This dual-model strategy highlights both the overwhelming predictive power of baseline glucose and the independent contributions of metabolic, inflammatory, and lipid-related factors.

To provide further interpretability, SHAP was applied to the random forest model in Model 2. SHAP enabled not only global ranking of variable importance but also visualization of the direction and magnitude of feature effects. Although the order of features sometimes differed from that in the averaged ML rankings, this reflects methodological differences rather than inconsistency. SHAP uniquely illustrates both feature contribution and directional effect, complementing conventional importance metrics.

Beyond regression, we also developed a binary classification model to predict prediabetes status at follow-up. This approach aligns with clinical decision making, which often relies on categorical thresholds (e.g., FPG ≥ 100 mg/dL) rather than gradual changes [7]. Our classifier achieved exceptionally high sensitivity (0.995), correctly identifying nearly all prediabetes cases. While specificity was low (0.010), this trade-off suits screening programs, where minimizing false negatives is crucial. The model's PR-AUC (0.873) confirmed strong precision–recall balance, especially important in moderately imbalanced datasets where ROC-AUC can be less informative [42].

The integration of modern ML algorithms—RF, SGB, XGBoost, and EN—into both regression and classification frameworks provides robust evidence. XGBoost and RF,

in particular, consistently demonstrated high performance and stability, underscoring their suitability for structured biomedical data. Compared with deep learning, which often requires very large datasets and provides limited interpretability, these tree-based methods are efficient and clinically transparent. By combining accuracy with SHAP-based interpretability, our framework identifies modifiable risk factors that can guide preventive strategies.

Importantly, predictors overlapped across continuous and binary frameworks. BF, TG, LDL-C, and age consistently emerged as influential, underscoring their central role in early glycemic risk. The binary model further revealed threshold effects, suggesting that some predictors disproportionately influence crossing into the prediabetic range. Thus, continuous models are better suited to track metabolic trajectories, whereas binary models directly support clinical screening. Together, they provide complementary insights into early dysglycemia risk.

The strong predictive role of FPG_{base} aligns with previous longitudinal studies such as the Bogalusa Heart Study [43], which showed that childhood FPG predicted adult dysglycemia over 21 years. This reinforces FPG as a simple yet powerful marker of glucose metabolism. Similarly, BF has long been linked with impaired glucose regulation. Jo et al. [44], using NHANES data, reported higher prevalence of abnormal glucose in normal-weight individuals with elevated BF compared with overweight individuals with lower BF. Our findings further support adiposity as a central risk factor, even in young men.

Chronic inflammation, reflected by higher WBC counts, is another recognized contributor to insulin resistance [45,46]. Jiang et al. [45] found that WBC was significantly higher in Chinese individuals with prediabetes, and similar results were observed in Japanese cohorts. While the relationship between WBC and β -cell function remains debated, obesity may partly explain these conflicting findings, as it is associated with both higher WBC and altered insulin secretion.

Age also played a role, despite our cohort being restricted to young adults. Aging is associated with deterioration in glucose metabolism through mechanisms including reduced β -cell mass, oxidative stress, and impaired DNA repair [47–54]. Even subtle age effects were detectable in this population, highlighting their early contribution to glycemic risk.

TSH levels also emerged as an important factor. Prior studies have suggested associations between thyroid dysfunction, adiposity, and glucose metabolism [55,56]. Our findings add to this evidence, suggesting that thyroid function may contribute to prediabetes risk, particularly in younger individuals. Finally, dyslipidemia, reflected by elevated TG and LDL-C, was positively associated with δ -FPG. These changes are hallmarks of insulin resistance [57,58] and further underline the interdependence between lipid and glucose metabolism.

Despite these strengths, some limitations must be acknowledged. First, plasma insulin and HbA1c were not routinely available in our dataset. These biomarkers are critical for assessing insulin resistance and β -cell function, and their absence limited the comprehensiveness of our models. Second, unmeasured lifestyle factors such as detailed dietary intake, sleep quality, and family history could not be fully accounted for, potentially introducing residual confounding. Third, our cohort consisted solely of young Taiwanese men. While this enhances internal validity, it limits generalizability to women, older adults, or other ethnicities. External validation in more diverse populations is needed. Finally, although we employed repeated cross-validation and confidence interval estimation to strengthen internal validity, external or temporal validation remains a priority for future research.

5. Conclusions

FPG_{base} was the strongest predictor of long-term glycemic change in young Taiwanese men. Excluding FPG_{base}, BF, WBC, age, TSH, LDL-C, and HDL-C remained key determinants. The dual-framework design—continuous regression and binary classification—provides complementary perspectives. Regression models capture subtle glycemic trajectories, while classification models enable threshold-based screening. Together, they enhance risk stratification and support targeted preventive strategies in young adults.

Supplementary Materials: The following supporting information can be downloaded at <https://www.mdpi.com/article/10.3390/diagnostics15192507/s1>. Figure S1: Feature correlation heatmap; Figure S2: Simplified version of our modeling pipeline; Table S1: Classification performance with 95% bootstrap confidence intervals for ROC_AUC and PR_AUC; Table S2: feature_vif_table.csv.

Author Contributions: Conceptualization, C.-C.Y.; software, T.-W.C.; validation, T.-W.C.; formal analysis, T.-W.C.; investigation, C.-C.Y. and T.-W.C.; data curation, S.-T.W.; writing—original draft preparation, C.-C.Y.; writing—review and editing, C.-H.L. and Y.-J.C. All authors have read and agreed to the published version of the manuscript.

Funding: This research was funded by the Kaohsiung Armed Forces General Hospital (grant number: KAFGH_E_112053).

Institutional Review Board Statement: The study was conducted in accordance with the Declaration of Helsinki and approved by the Institutional Review Board of the Kaohsiung Armed Forces General Hospital 1 (protocol code KAFGHIRB 112-006, date of approval 21 June 2023).

Informed Consent Statement: Not applicable.

Data Availability Statement: Data available on request due to privacy/ethical restrictions. A minimal, anonymized analysis code template is available from the corresponding author upon reasonable request.

Acknowledgments: The authors thank all subjects who participated in the study.

Conflicts of Interest: The authors declare no conflicts of interest.

References

1. Ong, K.L.; Stafford, L.K.; McLaughlin, S.A.; Boyko, E.J.; Vollset, S.E.; Smith, A.E.; Dalton, B.E.; Duprey, J.; Cruz, J.A.; Hagins, H.; et al. Global, regional, and national burden of diabetes from 1990 to 2021, with projections of prevalence to 2050: A systematic analysis for the Global Burden of Disease Study 2021. *Lancet* **2023**, *402*, 203–234. [[CrossRef](#)]
2. International Diabetes Federation. *IDF Diabetes Atlas*, 10th ed.; International Diabetes Federation: Brussels, Belgium, 2021. Available online: <https://www.diabetesatlas.org> (accessed on 6 December 2021).
3. Hou, Y.C.; Feng, H.C.; Tzeng, I.S.; Kuo, C.Y.; Cheng, C.F.; Wu, J.H.; Yang, S.H. Dietary patterns and the risk of prediabetes in Taiwan: A cross-sectional study. *Nutrients* **2020**, *12*, 3322. [[CrossRef](#)]
4. Yu, N.C.; Su, H.Y.; Tsai, S.T.; Lin, B.J.; Shiu, R.-S.; Hsieh, Y.-C.; Sheu, W.H.-H. ABC control of diabetes: Survey data from National Diabetes Health Promotion Centers in Taiwan. *Diabetes Res. Clin. Pract.* **2009**, *84*, 194–200. [[CrossRef](#)]
5. Patel, B.; Patel, C.; Panchal, D.; Patel, S. A retrospective evaluation of the trend of prevalence of type 2 diabetes mellitus in different age groups in a tertiary care hospital. *Panacea J. Med. Sci.* **2021**, *11*, 130–133. [[CrossRef](#)]
6. Zoungas, S.; Woodward, M.; Li, Q.; Cooper, M.E.; Hamet, P.; Harrap, S.; Heller, S.; Marre, M.; Patel, A.; Poulter, N.; et al. Impact of age, age at diagnosis and duration of diabetes on the risk of macrovascular and microvascular complications and death in type 2 diabetes. *Diabetologia* **2014**, *57*, 2465–2474. [[CrossRef](#)] [[PubMed](#)]
7. American Diabetes Association Professional Practice Committee. 2. Diagnosis and Classification of Diabetes: Standards of Care in Diabetes—2025. *Diabetes Care* **2025**, *48* (Suppl. 1), S27–S49. [[CrossRef](#)]
8. Huang, X.; Han, Y.; Jang, K.; Kim, M. Early Prediction for Prediabetes and Type 2 Diabetes Using the Genetic Risk Score and Oxidative Stress Score. *Antioxidants* **2022**, *11*, 1196. [[CrossRef](#)]
9. Wu, J.; Zhou, J.; Yin, X.; Chen, Y.; Lin, X.; Xu, Z.; Li, H. A Prediction Model for Prediabetes Risk in Middle-Aged and Elderly Populations: A Prospective Cohort Study in China. *Int. J. Endocrinol.* **2021**, *2021*, 2520806. [[CrossRef](#)]

10. Yuk, H.; Gim, J.; Min, J.K.; Yun, J.; Heo, T.Y. Artificial intelligence-based prediction of diabetes and prediabetes using health checkup data in Korea. *Appl. Artif. Intell.* **2022**, *36*, 2145644. [[CrossRef](#)]
11. Wu, C.Z.; Huang, L.Y.; Chen, F.Y.; Kuo, C.H.; Yeih, D.F. Using Machine Learning to Predict Abnormal Carotid Intima-Media Thickness in Type 2 Diabetes. *Diagnostics* **2023**, *13*, 1834. [[CrossRef](#)]
12. Liu, C.H.; Chang, C.F.; Chen, I.C.; Lin, F.-M.; Tzou, S.-J.; Hsieh, C.-B.; Chu, T.-W.; Pei, D. Machine Learning Prediction of Prediabetes in a Young Male Chinese Cohort with 5.8-Year Follow-Up. *Diagnostics* **2024**, *14*, 979. [[CrossRef](#)]
13. Lundberg, S.M.; Lee, S.I. A unified approach to interpreting model predictions. In Proceedings of the Advances in neural information processing systems, Long Beach, CA, USA, 4–9 December 2017; Volume 30, pp. 4765–4774.
14. Lundberg, S.M.; Erion, G.; Chen, H.; DeGrave, A.; Prutkin, J.M.; Nair, B.; Katz, R.; Himmelfarb, J.; Bansal, N.; Lee, S.I. From local explanations to global understanding with explainable AI for trees. *Nat. Mach. Intell.* **2020**, *2*, 56–67. [[CrossRef](#)]
15. Ponce-Bobadilla, A.V.; Schmitt, V.; Maier, C.S.; Mensing, S.; Stodtmann, S. Practical guide to SHAP analysis: Explaining supervised machine learning model predictions in drug development. *Clin. Transl. Sci.* **2024**, *17*, e70056. [[CrossRef](#)] [[PubMed](#)]
16. Sadeghi, Z.; Alizadehsani, R.; Cifci, M.A.; Kausar, S.; Rehman, R.; Mahanta, P.; Bora, P.K.; Almasri, A.; Alkhawaldeh, R.S.; Hussain, S.; et al. A review of Explainable Artificial Intelligence in healthcare. *Comput. Electr. Eng.* **2024**, *118*, 109370. [[CrossRef](#)]
17. Deo, R.C. Machine learning in medicine. *Circulation* **2015**, *132*, 1920–1930. [[CrossRef](#)] [[PubMed](#)]
18. Beam, A.L.; Kohane, I.S. Big data and machine learning in health care. *JAMA* **2018**, *319*, 1317–1318. [[CrossRef](#)]
19. Sidey-Gibbons, J.A.; Sidey-Gibbons, C.J. Machine learning in medicine: A practical introduction. *BMC Med. Res. Methodol.* **2019**, *19*, 64. [[CrossRef](#)]
20. Wu, X.; Tsai, S.P.; Tsao, C.K.; Chiu, M.L.; Tsai, M.K.; Lu, P.J.; Lee, J.H.; Chen, C.H.; Wen, C.; Chang, S.-S.; et al. Cohort Profile: The Taiwan MJ Cohort: Half a million Chinese with repeated health surveillance data. *Int. J. Epidemiol.* **2017**, *46*, 1744–1744g. [[CrossRef](#)]
21. MJ Health Research Foundatio. The Introduction of MJ Health Database. MJ Health Research Foundation Technical Report, MJHRF-TR-01. 2016. Available online: <http://www.mjhrf.org/upload/user/files/MJHRF-TR-01%20MJ%20Health%20Database.pdf> (accessed on 22 August 2016).
22. Breiman, L. Random forests. *Mach. Learn.* **2001**, *45*, 5–32. [[CrossRef](#)]
23. Friedman, J.H. Greedy function approximation: A gradient boosting machine. *Ann. Stat.* **2001**, *29*, 1189–1232. [[CrossRef](#)]
24. Ester, M.; Kriegel, H.P.; Xu, X.J.G.A. XGBoost: A scalable tree boosting system. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, 13–17 August 2016; Geographical Analysis. p. 785.
25. Zou, H.; Hastie, T. Regularization and variable selection via the elastic net. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **2005**, *67*, 301–320. [[CrossRef](#)]
26. Molnar, C. *Interpretable Machine Learning*; Lulu.com: Durham, NC, USA, 2020; ISBN 0244768528/9780244768522. Available online: <https://christophm.github.io/interpretable-ml-book/> (accessed on 22 August 2016).
27. Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; et al. Scikit-learn: Machine learning in Python. *J. Mach. Learn. Res.* **2011**, *12*, 2825–2830.
28. Chawla, N.V.; Bowyer, K.W.; Hall, L.O.; Kegelmeyer, W.P. SMOTE: Synthetic minority over-sampling technique. *J. Artif. Intell. Res.* **2002**, *16*, 321–357. [[CrossRef](#)]
29. Batista, G.E.; Prati, R.C.; Monard, M.C. A study of the behavior of several methods for balancing machine learning training data. *ACM SIGKDD Explor. Newsl.* **2004**, *6*, 20–29. [[CrossRef](#)]
30. Salimi, M.; Atif, D.; Oliva, D.; Abraham, A.; Ventura, S. Handling imbalanced medical datasets: Review of a decade of research. *Artif. Intell. Rev.* **2024**, *57*, 273. [[CrossRef](#)]
31. Gurcan, F.; Soylu, A. Learning from imbalanced data: Integration of advanced resampling techniques and machine learning models for enhanced cancer diagnosis and prognosis. *Cancers* **2024**, *16*, 3417. [[CrossRef](#)] [[PubMed](#)]
32. Akiba, T.; Sano, S.; Yanase, T.; Ohta, T.; Koyama, M. Optuna: A next-generation hyperparameter optimization framework. In Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, Anchorage, AK, USA, 4–8 August 2019; pp. 2623–2631.
33. Saito, T.; Rehmsmeier, M. The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. *PLoS ONE* **2015**, *10*, e0118432. [[CrossRef](#)]
34. Collins, G.S.; Dhiman, P.; Ma, J.; Schlussel, M.M.; Archer, L.; Van Calster, B.; Harrell, F.E.; Martin, G.P.; Moons, K.G.M.; van Smeden, M.; et al. Evaluation of clinical prediction models (part 1): From development to external validation. *BMJ* **2024**, *384*, e074819. [[CrossRef](#)]
35. Riley, R.D.; Archer, L.; Snell, K.I.; Ensor, J.; Dhiman, P.; Martin, G.P.; Bonnett, L.J.; Collins, G.S. Evaluation of clinical prediction models (part 2): How to undertake an external validation study. *BMJ* **2024**, *384*, e074820. [[CrossRef](#)]
36. Vickers, A.J.; Elkin, E.B. Decision curve analysis: A novel method for evaluating prediction models. *Med. Decis. Mak.* **2006**, *26*, 565–574. [[CrossRef](#)]

37. Vickers, A.J.; van Calster, B.; Steyerberg, E.W. A simple, step-by-step guide to interpreting decision curve analysis. *Diagn. Progn. Res.* **2019**, *3*, 18. [[CrossRef](#)]
38. Student. The probable error of a mean. *Biometrika* **1908**, *6*, 1–25. [[CrossRef](#)]
39. Altman, D.G.; Bland, J.M. Standard deviations and standard errors. *BMJ* **2005**, *331*, 903. [[CrossRef](#)]
40. Cumming, G.; Finch, S. Inference by eye: Confidence intervals and how to read pictures of data. *Am. Psychol.* **2005**, *60*, 170. [[CrossRef](#)]
41. Dormann, C.F.; Elith, J.; Bacher, S.; Buchmann, C.; Carl, G.; Carré, G.; Marquéz, J.R.G.; Gruber, B.; Lafourcade, B.; Leitão, P.J.; et al. Collinearity: A review of methods to deal with it and a simulation study evaluating their performance. *Ecography* **2013**, *36*, 27–46. [[CrossRef](#)]
42. Davis, J.; Goadrich, M. The relationship between Precision-Recall and ROC curves. In Proceedings of the 23rd international conference on Machine learning, Pittsburgh, PA, USA, 25–29 June 2006; pp. 233–240.
43. Eraso, L.H.; Ginwala, N.; Qasim, A.N.; Mehta, N.N.; Dlugash, R.; Kapoor, S.; Schwartz, S.; Schutta, M.; Iqbal, N.; Mohler, E.R.; et al. Association of lower plasma fetuin-a levels with peripheral arterial disease in type 2 diabetes. *Diabetes Care* **2010**, *33*, 408–410. [[CrossRef](#)] [[PubMed](#)]
44. Jo, A.; Mainous, A.G., 3rd. Informational value of percent body fat with body mass index for the risk of abnormal blood glucose: A nationally representative cross-sectional study. *BMJ Open* **2018**, *8*, e019200. [[CrossRef](#)] [[PubMed](#)]
45. Jiang, H.; Yan, W.H.; Li, C.J.; Wang, A.P.; Dou, J.T.; Mu, Y.M. Elevated white blood cell count is associated with higher risk of glucose metabolism disorders in middle-aged and elderly Chinese people. *Int. J. Env. Environ. Res. Public. Health.* **2014**, *11*, 5497–5509. [[CrossRef](#)]
46. Ohshita, K.; Yamane, K.; Hanafusa, M.; Mori, H.; Mito, K.; Okubo, M.; Hara, H.; Kohno, N. Elevated white blood cell count in subjects with impaired glucose tolerance. *Diabetes Care* **2004**, *27*, 491–496. [[CrossRef](#)]
47. Chen, L.K.; Lin, M.H.; Chen, Z.J.; Hwang, S.J.; Chiou, S.T. Association of insulin resistance and hematologic parameters: Study of a middle-aged and elderly Chinese population in Taiwan. *J. Chin. Med. Assoc.* **2006**, *69*, 248–253. [[CrossRef](#)]
48. Boemi, M.; Furlan, G.; Luconi, M.P. Chapter 29—Glucose Metabolism, Insulin, and Aging: Role of Nutrition. In *Molecular Basis of Nutrition and Aging*; Malavolta, M., Moccagiani, E., Eds.; Academic Press: Cambridge, MA, USA, 2016; pp. 393–409.
49. U.S. Department of Health and Human Services. Estimates of diabetes its burden in the United States 2014 Atlanta (GA). In *Centers for Disease Control Prevention National Diabetes Statistics Report*; U.S. Department of Health and Human Services: Washington, DC, USA, 2014.
50. Jacob, K.D.; Noren Hooten, N.; Trzeciak, A.R.; Evans, M.K. Markers of oxidant stress that are clinically relevant in aging and age-related disease. *Mech. Ageing Dev.* **2013**, *134*, 139–157. [[CrossRef](#)]
51. Lombard, D.B.; Chua, K.F.; Mostoslavsky, R.; Franco, S.; Gostissa, M.; Alt, F.W. DNA repair, genome stability, and aging. *Cell* **2005**, *120*, 497–512. [[CrossRef](#)] [[PubMed](#)]
52. Haus, J.M.; Carrithers, J.A.; Trappe, S.W.; Trappe, T.A. Collagen, cross-linking, and advanced glycation end products in aging human skeletal muscle. *J. Appl. Physiol.* **2007**, *103*, 2068–2076. [[CrossRef](#)]
53. Paneni, F.; Costantino, S.; Cosentino, F. Molecular pathways of arterial aging. *Clin. Sci.* **2015**, *128*, 69–79. [[CrossRef](#)] [[PubMed](#)]
54. Røder, M.E.; Schwartz, R.S.; Prigeon, R.L.; Kahn, S.E. Reduced pancreatic B cell compensation to the insulin resistance of aging: Impact on proinsulin and insulin levels. *J. Clin. Endocrinol. Metab.* **2000**, *85*, 2275–2280. [[CrossRef](#)]
55. Korzeniowska, K.A.; Brzeziński, M.; Szarejko, K.; Radziwiłł, M.; Anyszek, T.; Czupryniak, L.; Soszyński, P.; Berggren, P.O.; Mysliwiec, M. The association of thyroid-stimulating hormone (TSH) and free thyroxine (fT4) concentration levels with carbohydrate and lipid metabolism in obese and overweight teenagers. *Endokrynol. Pol.* **2019**, *70*, 172–178. [[CrossRef](#)]
56. Kim, H.; Jung, D.Y.; Lee, S.H.; Cho, J.H.; Yim, H.W.; Kim, H.S. Retrospective cohort analysis comparing changes in blood glucose level and body composition according to changes in thyroid-stimulating hormone level. *J. Diabetes.* **2022**, *14*, 620–629. [[CrossRef](#)]
57. Cohn, G.; Valdes, G.; Capuzzi, D.M. Pathophysiology and treatment of the dyslipidemia of insulin resistance. *Curr. Cardiol. Rep.* **2001**, *3*, 416–423. [[CrossRef](#)]
58. Savage, D.B.; Petersen, K.F.; Shulman, G.I. Mechanisms of insulin resistance in humans and possible links with inflammation. *Hypertension* **2005**, *45*, 828–833. [[CrossRef](#)]

Disclaimer/Publisher’s Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.